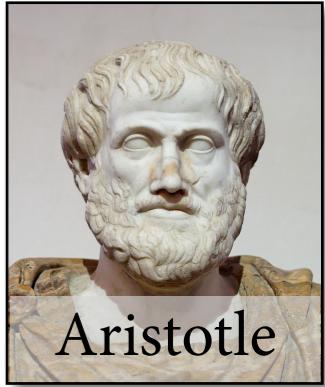


ANLP

18 - semantics

David Schlangen
University of Potsdam, MSc Cognitive Systems
Winter 2019 / 2020

Computing with meanings



- Ancient problem: *inference*.
 - ▶ How can we tell whether a sentence follows from others?
 - ▶ Can we compute this automatically?

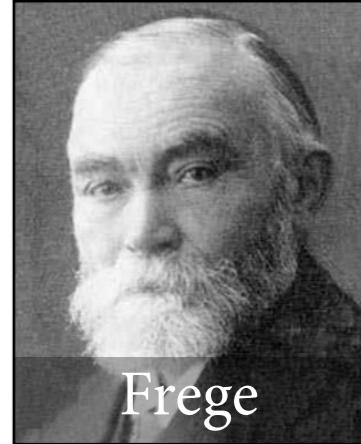
All men are mortal.

Socrates is a man.

Therefore, Socrates is mortal.

Formal meaning representations

- Aristotle with more modern tools (ca. 2000):
 - ▶ Compute *meaning representation* in some formal language (e.g. predicate logic)
 - ▶ so that it captures something relevant about the sentence's meaning (e.g. its *truth conditions*)
 - ▶ and then use reasoning tools for the formal language (e.g. a *theorem prover* for predicate logic)



All men are mortal.

Socrates is a man.

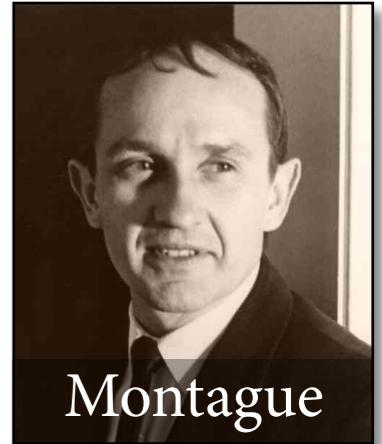
Therefore, Socrates is mortal.

$\forall x. \text{man}(x) \rightarrow \text{mortal}(x)$

$\text{man}(s)$

$\text{mortal}(s)$

Compositional semantics



$S \rightarrow NP\ VP$

$\langle S \rangle = \langle NP \rangle (\langle VP \rangle)$

$VP \rightarrow V\ NP$

$\langle VP \rangle = \lambda y \langle NP \rangle (\langle V \rangle (y))$

$NP \rightarrow Det\ N$

$\langle NP \rangle = \langle Det \rangle (\langle N \rangle)$

$NP \rightarrow John$

$\langle NP \rangle = \lambda P\ P(j^*)$

$V \rightarrow eats$

$\langle V \rangle = eat'$

$Det \rightarrow a$

$\langle Det \rangle = \lambda P \lambda Q \exists x P(x) \wedge Q(x)$

$N \rightarrow sandwich$

$\langle N \rangle = sw'$

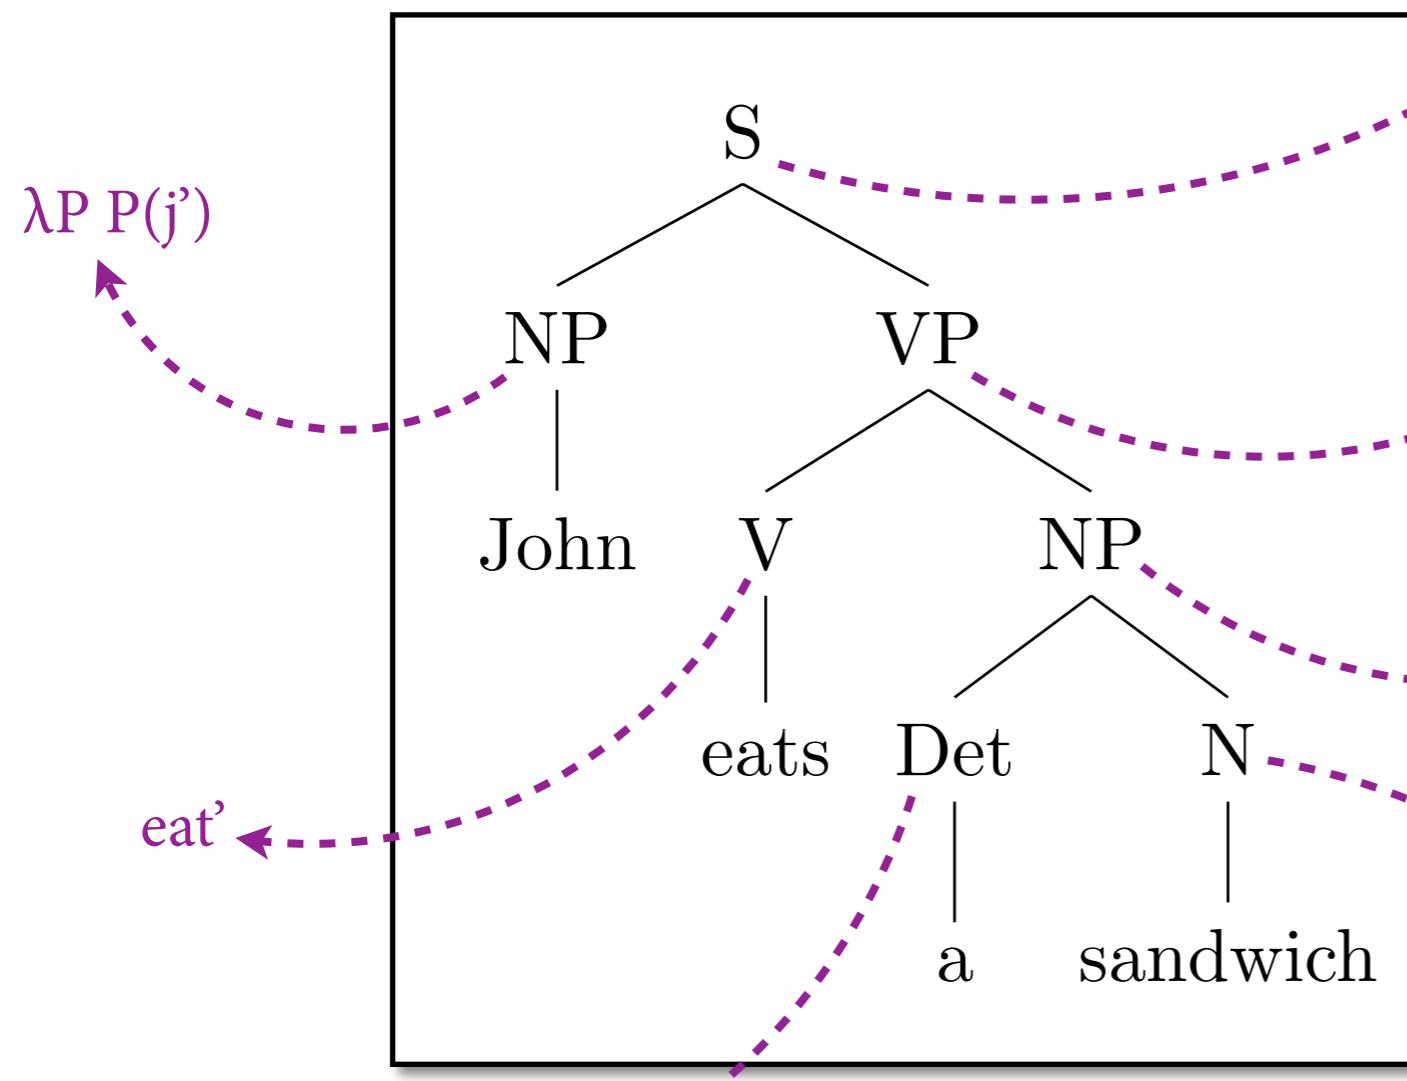


when you apply this
syntax rule ...



... construct λ -term for parent
from λ -terms for children like this

Example



$\lambda P \lambda Q \exists x P(x) \wedge Q(x)$

$$\begin{aligned}
 & (\lambda P P(j')) (\lambda y \exists x sw'(x) \wedge eat'(y)(x)) \\
 \rightarrow_{\beta} & (\lambda y \exists x sw'(x) \wedge eat'(y)(x))(j') \\
 \rightarrow_{\beta} & \exists x sw'(x) \wedge eat'(j')(x)
 \end{aligned}$$

$$\begin{aligned}
 & \lambda y (\lambda Q \exists x sw'(x) \wedge Q(x))(eat'(y)) \\
 \rightarrow_{\beta} & \lambda y \exists x sw'(x) \wedge eat'(y)(x)
 \end{aligned}$$

$$\begin{aligned}
 & (\lambda P \lambda Q \exists x P(x) \wedge Q(x))(sw') \\
 \rightarrow_{\beta} & \lambda Q \exists x sw'(x) \wedge Q(x)
 \end{aligned}$$

sw'

Semantic parsing

- Open issue in classical semantics construction:
Where do we get large grammar that supports it?
- Current trend in CL is *semantic parsing*:
learn mapping from sentence to formal meaning representation using statistical methods.
- E.g. from Geoquery corpus (880 sentences):

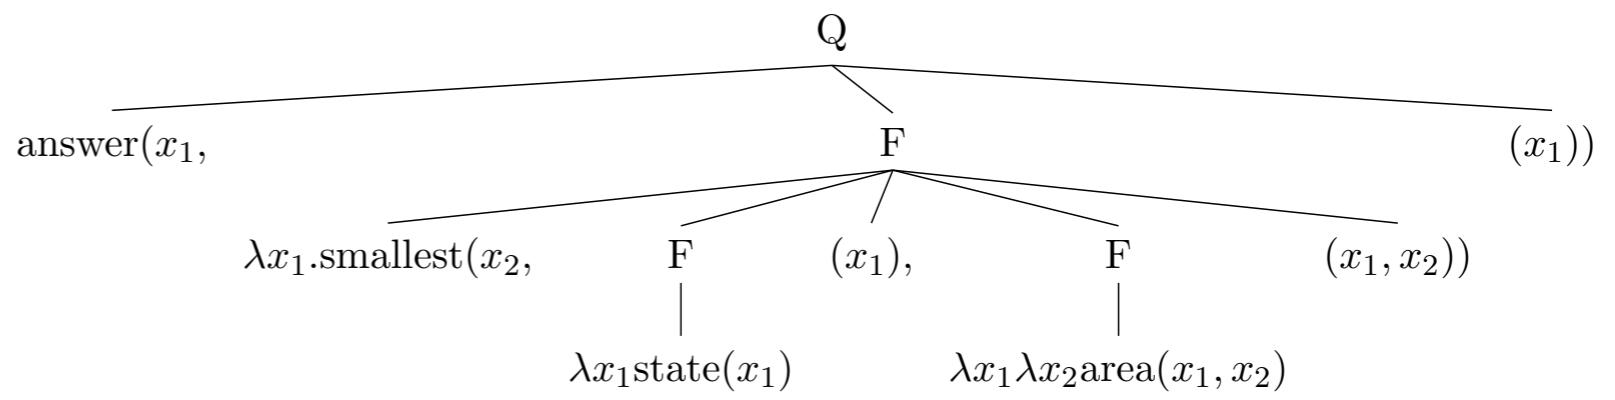
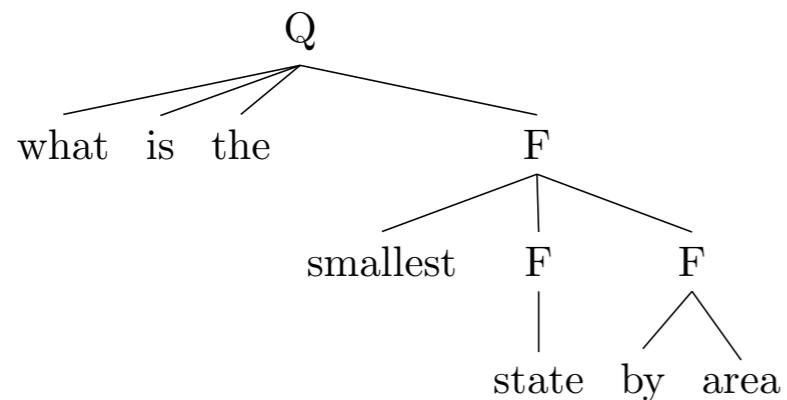
What is the smallest state by area?

```
answer(x1, smallest(x2, state(x1), area(x1, x2)))
```

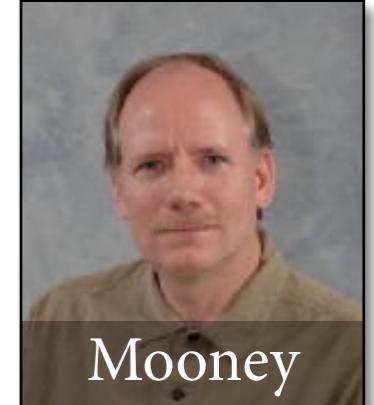
With synchronous grammars

- Use a synchronous grammar (\approx SCFG) to simultaneously generate strings and λ -expressions.

$Q \rightarrow \text{what is the } F$	$Q \rightarrow \text{answer}(x_1, F(x_1))$
$F \rightarrow \text{smallest } F \ F$	$F \rightarrow \lambda x_1 \text{ smallest}(x_2, F(x_1), F(x_1, x_2))$
$F \rightarrow \text{state}$	$F \rightarrow \lambda x_1 \text{ state}(x_1)$
$F \rightarrow \text{by area}$	$F \rightarrow \lambda x_1 \lambda x_2 \text{ area}(x_1, x_2)$



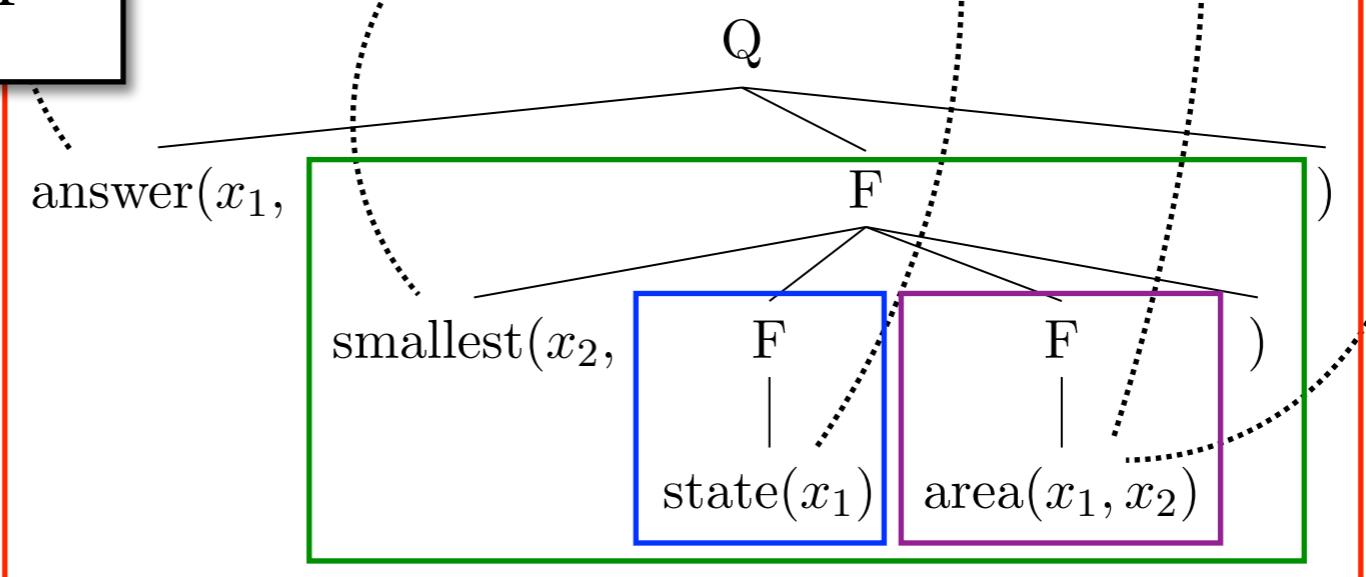
Wong & Mooney



what is the smallest state by area

Where do unaligned words belong?

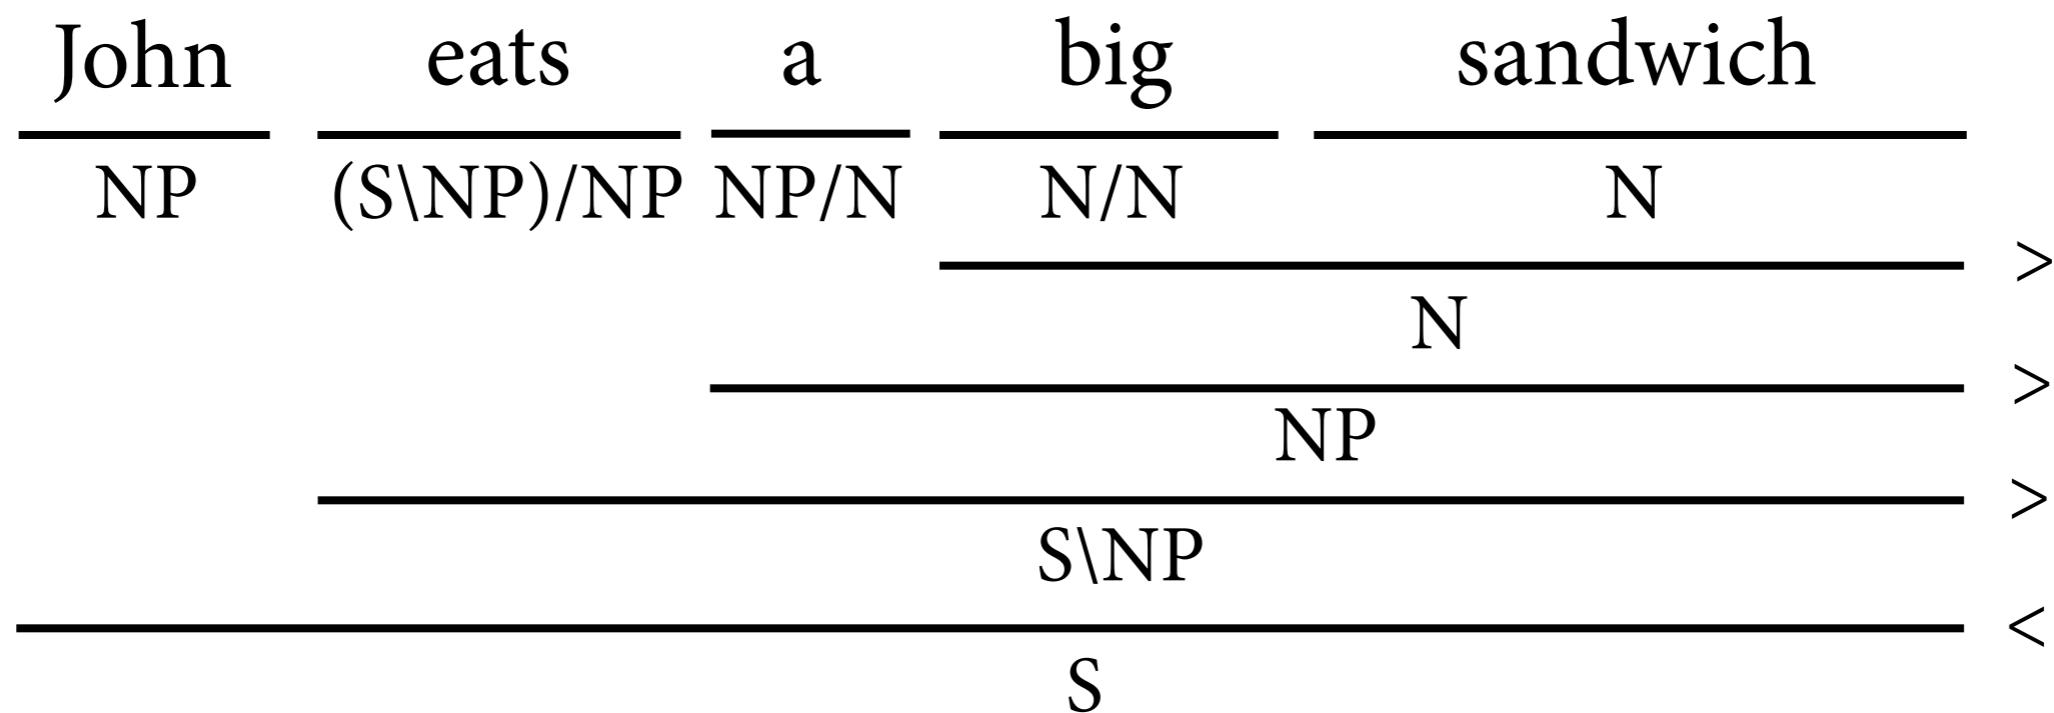
$Q \rightarrow \text{what is the F} \mid F \rightarrow \text{smallest F}$
 $Q \rightarrow \text{what F} \mid F \rightarrow \text{is the smallest F}$



Assumptions:

- alignments between words and nodes
- unambiguous structure of meaning representation

Combinatory categorial grammar



Semantics in CCG

$$\frac{X: a}{Y/(Y\backslash X): \lambda P.P(a)} >T$$

$$\frac{X/Y: f \quad Y/Z: g}{X/Z: \lambda x.f(g(x))} >B$$

$$\frac{X/Y: f \quad Y\backslash Z: g}{X\backslash Z: \lambda x.f(g(x))} >Bx$$

$$\frac{X: a}{Y\backslash(Y/X): \lambda P.P(a)} <T$$

$$\frac{Y\backslash Z: g \quad X\backslash Y: f}{X\backslash Z: \lambda x.f(g(x))} <B$$

$$\frac{Y/Z: g \quad X\backslash Y: f}{X/Z: \lambda x.f(g(x))} <Bx$$

$$\begin{array}{c}
 \text{John} \\
 \hline
 \text{NP: } h^* \\
 \hline
 S/(S\backslash NP): \lambda P.P(h^*) >T \\
 \hline
 S/NP: \lambda x.(\lambda P.P(h^*))(eat'(x)) \Rightarrow_{\beta} \lambda x.eat'(x)(h^*) \\
 \hline
 S: (\lambda x.eat'(x)(h^*))(sw') \Rightarrow_{\beta} eat'(sw')(h^*)
 \end{array}$$

eats

$$\begin{array}{c}
 \text{a sandwich} \\
 \hline
 \text{NP: } sw' \\
 \hline
 \end{array}
 >$$

Zettlemoyer & Collins

GENLEX: build candidates for lexicon entries

Rules		Categories produced from logical form
Input Trigger	Output Category	$\arg \max(\lambda x. state(x) \wedge borders(x, \text{texas}), \lambda x. size(x))$
constant c	$NP : c$	$NP : \text{texas}$
arity one predicate p_1	$N : \lambda x. p_1(x)$	$N : \lambda x. state(x)$
arity one predicate p_1	$S \setminus NP : \lambda x. p_1(x)$	$S \setminus NP : \lambda x. state(x)$
arity two predicate p_2	$(S \setminus NP) / NP : \lambda x. \lambda y. p_2(y, x)$	$(S \setminus NP) / NP : \lambda x. \lambda y. borders(y, x)$
arity two predicate p_2	$(S \setminus NP) / NP : \lambda x. \lambda y. p_2(x, y)$	$(S \setminus NP) / NP : \lambda x. \lambda y. borders(x, y)$
arity one predicate p_1	$N / N : \lambda g. \lambda x. p_1(x) \wedge g(x)$	$N / N : \lambda g. \lambda x. state(x) \wedge g(x)$
literal with arity two predicate p_2 and constant second argument c	$N / N : \lambda g. \lambda x. p_2(x, c) \wedge g(x)$	$N / N : \lambda g. \lambda x. borders(x, \text{texas}) \wedge g(x)$
arity two predicate p_2	$(N \setminus N) / NP : \lambda x. \lambda g. \lambda y. p_2(x, y) \wedge g(x)$	$(N \setminus N) / NP : \lambda g. \lambda x. \lambda y. borders(x, y) \wedge g(x)$
an arg max / min with second argument arity one function f	$NP / N : \lambda g. \arg \max / \min(g, \lambda x. f(x))$	$NP / N : \lambda g. \arg \max(g, \lambda x. size(x))$
an arity one numeric-ranged function f	$S / NP : \lambda x. f(x)$	$S / NP : \lambda x. size(x)$

Log-linear probability models

- Define probability of parse tree in terms of *features*:

$$P(t \mid w) = \frac{e^{\theta \cdot f(t, w)}}{\sum_{t'} e^{\theta \cdot f(t', w)}}$$

where $\theta \cdot f(t, w) = \theta_1 \cdot f_1(t, w) + \dots + \theta_n \cdot f_n(t, w)$

- Features $f(t, w)$ can capture arbitrary properties of t and w .
 - ▶ Here: Each feature counts uses of one grammar rule.
- Train weight vector θ from data.

Zettlemoyer & Collins

overall learning algorithm

Algorithm:

- For $t = 1 \dots T$

Step 1: (Lexical generation)

- For $i = 1 \dots n$:
 - Set $\lambda = \Lambda_0 \cup \text{GENLEX}(S_i, L_i)$.
 - Calculate $\pi = \text{PARSE}(S_i, L_i, \lambda, \bar{\theta}^{t-1})$.
 - Define λ_i to be the set of lexical entries in π .
- Set $\Lambda_t = \Lambda_0 \cup \bigcup_{i=1}^n \lambda_i$

Step 2: (Parameter Estimation)

- Set $\bar{\theta}^t = \text{ESTIMATE}(\Lambda_t, E, \bar{\theta}^{t-1})$

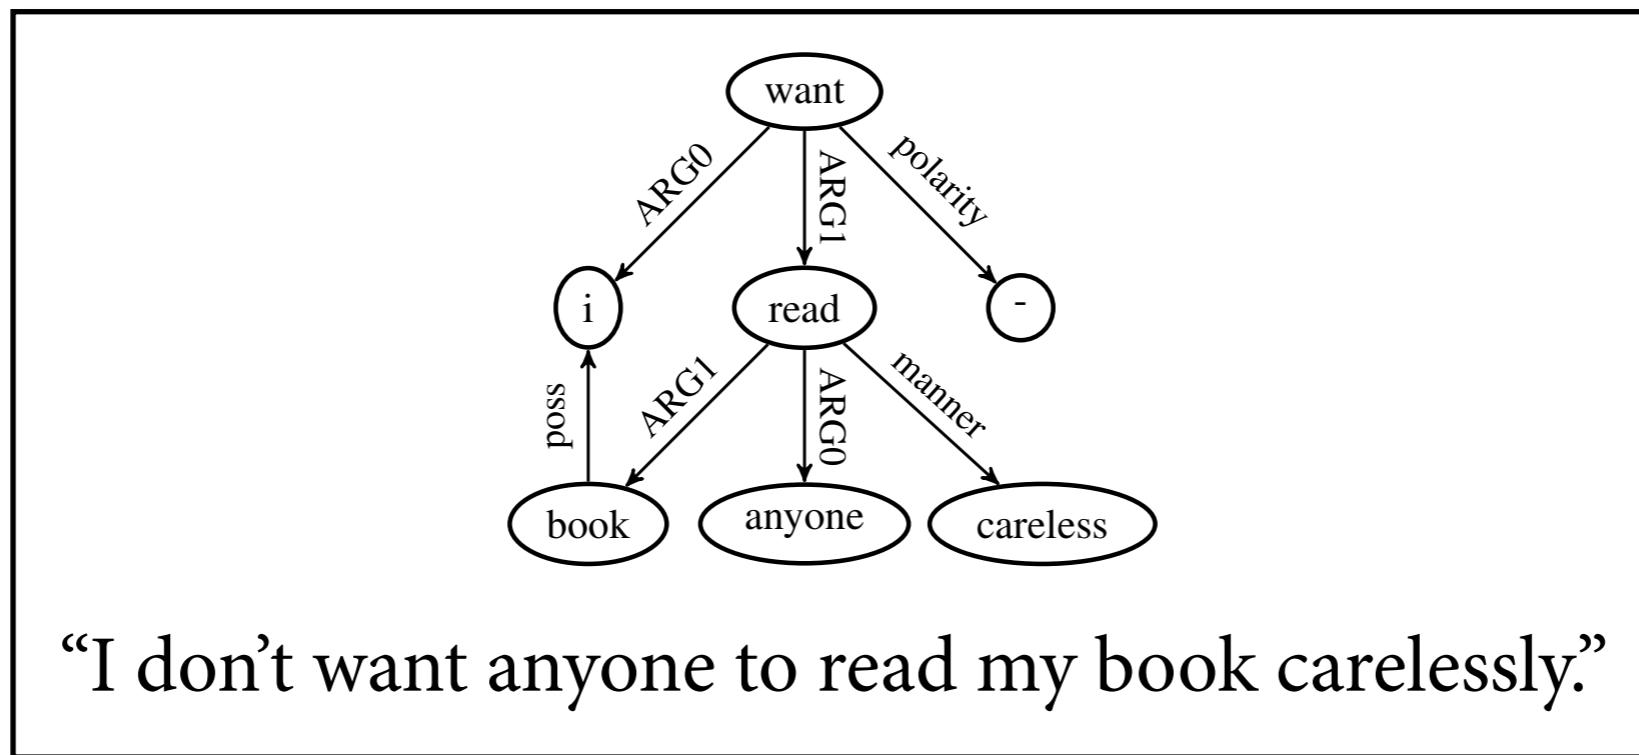
Evaluation results

System	Variable Free			Lambda Calculus		
	Rec.	Pre.	F1	Rec.	Pre.	F1
Cross Validation Results						
KRISP	71.7	93.3	81.1	–	–	–
WASP	74.8	87.2	80.5	–	–	–
Lu08	81.5	89.3	85.2	–	–	–
λ -WASP	–	–	–	86.6	92.0	89.2
Independent Test Set						
ZC05	–	–	–	79.3	96.3	87.0
ZC07	–	–	–	86.1	91.6	88.8
UBL	81.4	89.4	85.2	85.0	94.1	89.3
UBL-s	84.3	85.2	84.7	87.9	88.5	88.2

(on Geoquery 880 corpus)

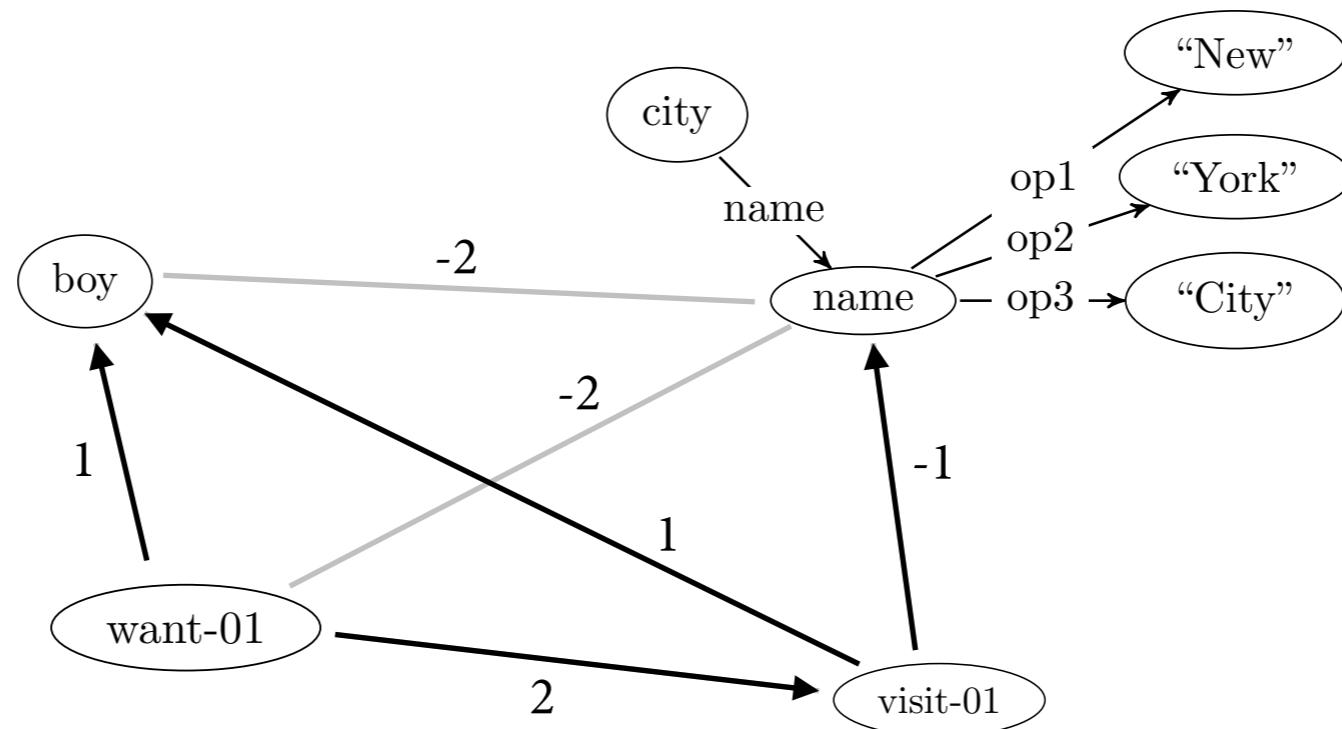
Abstract Meaning Representations

- Pros and cons of Geoquery:
 - ▶ semantic representations are trees — (too) simple
 - ▶ very small
- Since 2014, much larger corpora available:
~40k AMRs, graphs as semantic representations.



Dependency-style AMR parsing

“The boy wants to visit New York City.”



Concept Identification: determine atomic graph for each word.

Relation Identification: add all edges with positive weight; then repeatedly add least negative edge that connects subgraphs.

Issues with JAMR

- JAMR can draw edge between any two nodes; syntactic structure of sentence used only indirectly.
- Semantic representations for words don't know anything about their semantic arguments.
- Edges for control verbs added arbitrarily, not because linguistic principle of control discovered.
- No notion of compositionality!

Conclusion

- Challenge in compositional semantic construction:
 - Where do we get large-scale grammars?
- Semantic parsing: Learn such grammars from corpora with semantic annotations.
 - ▶ GeoQuery: small corpus of trees
 - ▶ AMRBank: new hotness
- Very active research topic right now.

slide credits

slides that look like this

Question 2: Tagging

- Given observations y_1, \dots, y_T , what is the most probable sequence x_1, \dots, x_T of hidden states?
- Maximum probability:
$$\max_{x_1, \dots, x_T} P(x_1, \dots, x_T | y_1, \dots, y_T)$$
- We are primarily interested in arg max:
$$\begin{aligned} & \arg \max_{x_1, \dots, x_T} P(x_1, \dots, x_T | y_1, \dots, y_T) \\ &= \arg \max_{x_1, \dots, x_T} \frac{P(x_1, \dots, x_T, y_1, \dots, y_T)}{P(y_1, \dots, y_T)} \\ &= \arg \max_{x_1, \dots, x_T} P(x_1, \dots, x_T, y_1, \dots, y_T) \end{aligned}$$

come from

earlier editions of this class (ANLP), given by Alexander Koller

and their use is gratefully acknowledged. I try to make any modifications obvious, but if there are errors on a slide, assume that I added them.